

# Proof-First Architectures for Independent Media Verification: Formal Requirements and Architectural Analysis

A Formal Systems Analysis of Trustless Integrity Verification  
Under Byzantine Adversarial Conditions

**Classification:** Distributed Systems / Applied Cryptography / Security

**Methodology:** Formal Analysis and Architectural Comparison

## Abstract

We formalize requirements for *independent media verification systems*—distributed infrastructures enabling cryptographic verification of digital artifact authenticity, temporal ordering, and authorship attribution by arbitrary third parties without reliance on trusted intermediaries. Under explicit Byzantine fault tolerance assumptions ( $f < n/3$  adversarial nodes), we derive verification complexity bounds for resource-constrained clients, construct a formal adversarial threat model, and demonstrate why conventional architectures fail to satisfy these requirements simultaneously.

Through formal complexity analysis, we establish that dual-ledger architectures with deterministic header commitments achieve  $O(1)$  verification complexity independent of global account population—a property we prove necessary for mass accessibility under resource constraints. We derive proof size bounds of 128–192 bytes and demonstrate verification costs of  $O(1)$  cryptographic operations plus  $O(|m|)$  hashing for artifact  $m$ . Comparative architectural analysis shows that header-first, account-isolated designs constitute the only known architecture class satisfying all specified constraints concurrently.

**Contributions:** (1) Formalized verification requirements with explicit resource bounds, (2) Byzantine adversarial model for media integrity systems, (3) Complexity proofs establishing browser-native feasibility, (4) Architectural necessity arguments for dual-ledger separation.

## 1. Introduction

The authenticity and provenance of digital media artifacts—documents, images, timestamped statements, recorded communications—increasingly determine critical outcomes in journalism, legal proceedings, scientific discourse, and institutional governance. Contemporary verification mechanisms require trust in centralized authorities: platform operators, certificate authorities, notary services, or institutional gatekeepers. This architectural dependency creates a structural vulnerability wherein degradation of trust in intermediaries—through technical compromise, institutional failure, or adversarial coercion—eliminates verification capability entirely.

This work addresses the following question: **Under what architectural constraints can media authenticity be verified independently of trusted third parties, with verification computationally feasible in resource-constrained environments, under Byzantine adversarial conditions?**

We establish formal requirements, derive necessary complexity bounds, and evaluate architectural compliance through comparative analysis. Our approach is descriptive—we characterize necessary properties without prescribing specific implementations, though we analyze reference architectures for concreteness.

### 1.1 Scope and Exclusions

This work concerns *integrity verification*—cryptographic proof that artifact  $m$  has not been modified since creation and existed at consensus time  $t$  with authorship attributable to public key  $pk$ . The following are explicitly outside our scope:

- Semantic correctness or truthfulness of content (requires external validation)
- Real-world identity binding beyond cryptographic key continuity
- Censorship resistance of content distribution (orthogonal to verification capability)
- Sybil resistance and reputation mechanisms (application-layer concerns)
- Privacy-preserving verification techniques (complementary but separate problem)

Our focus is establishing that artifact  $m$  existed in form  $H(m)$  at consensus time  $t$ , verifiable by any party possessing tuple  $(m, \pi, H)$  where  $\pi$  is a proof bundle and  $H$  represents consensus headers.

### 1.2 Terminology and Notation

We establish core terminology used throughout:

- **Artifact:** Digital content object subject to verification (document, image, statement), denoted  $m$
- **Commitment:** Cryptographic binding  $C = (h, t, pk)$  where  $h = H(m)$ ,  $t \in \mathbb{N}$  (timestamp),  $pk \in PK$  (public key space)
- **Proof bundle:** Data structure  $\pi$  enabling verification without external queries beyond peer-to-peer header retrieval

- **Verification predicate:** Boolean function  $\text{Verify}(m, \pi, H) \rightarrow \{\text{true}, \text{false}\}$  computable by any party
- **Header-first consensus:** Byzantine agreement mechanism finalizing commitments without executing state transitions
- **Account-local property:** Verification complexity independent of unrelated accounts

Throughout,  $N$  denotes total account count,  $n$  denotes consensus node count,  $f$  denotes Byzantine node count,  $\lambda$  denotes security parameter (typically 256), and  $\text{negl}(\lambda)$  denotes negligible probability in  $\lambda$ .

## 2. Adversarial Model and Trust Assumptions

We formalize the adversarial environment and underlying cryptographic assumptions.

### 2.1 Adversary Capabilities

We model a probabilistic polynomial-time adversary  $\mathbb{B}$  with the following capabilities:

- **Platform Control:**  $\mathbb{B}$  operates centralized infrastructure controlling artifact visibility, metadata, and archival access
- **Content Suppression:**  $\mathbb{B}$  removes artifacts from infrastructure under its operational control
- **Metadata Manipulation:**  $\mathbb{B}$  arbitrarily alters timestamps, authorship attribution, and ordering information on controlled platforms
- **Network Observation:**  $\mathbb{B}$  monitors network communications but cannot execute global network partition
- **Byzantine Node Control:**  $\mathbb{B}$  controls  $f < n/3$  consensus nodes exhibiting arbitrary Byzantine behavior
- **Equivocation Attempts:**  $\mathbb{B}$  presents conflicting artifact versions to different verifiers

### 2.2 Cryptographic Primitives

We assume standard cryptographic primitives with security parameter  $\lambda \geq 256$ :

**Hash Function Security.** Let  $H: \{0,1\}^* \rightarrow \{0,1\}^{\lambda}$  be a cryptographic hash function. We assume collision resistance: for any PPT adversary  $\mathbb{B}$ ,  $\Pr[\mathbb{B} \text{ outputs } x \neq y \text{ with } H(x) = H(y)] \leq \text{negl}(\lambda)$ .

**Digital Signature Security.** Let  $\Sigma = (\text{KeyGen}, \text{Sign}, \text{Verify})$  be a signature scheme. We assume existential unforgeability under adaptive chosen message attack (EUF-CMA): no PPT adversary can produce valid signature on message not previously signed, except with probability  $\text{negl}(\lambda)$ .

**Random Oracle Assumption.** Complexity analysis assumes hash functions behave as random oracles where applicable.

### 2.3 Byzantine Consensus Assumptions

We assume a Byzantine fault-tolerant consensus protocol with the following properties:

**Safety.** If  $f < n/3$  nodes exhibit Byzantine behavior, no two honest nodes finalize conflicting commitments for identical logical position.

**Liveness.** Under partial synchrony (GST model: unknown Global Stabilization Time after which network behaves synchronously), transactions submitted by honest nodes are eventually finalized.

**Header Availability.** Consensus headers achieve eventual consistency across honest nodes. We do not assume synchronous availability—verifiers may observe finalized headers with bounded but non-deterministic delay.

**Remark.** Temporary liveness failures delay verification but do not compromise safety. A cryptographically valid commitment remains valid indefinitely regardless of network conditions.

## 2.4 Excluded Attack Vectors

The following attack vectors lie outside our threat model:

- Permanent global network partition isolating all verifiers from consensus
- Total eclipse attacks compromising all peer connections of a verifier simultaneously
- Breaks in underlying cryptographic primitives (hash collisions, signature forgery)
- Quantum adversaries (post-quantum signature schemes are orthogonal extension)
- Implementation-specific side-channel attacks on client software

These represent either (1) fundamental failures affecting all cryptographic systems equally, or (2) implementation concerns orthogonal to architectural requirements.

## 2.5 Verification Objectives

Given artifact  $m$ , proof bundle  $\pi$ , and consensus headers  $H$ , verification establishes:

1. **Integrity:**  $H(m) = h$  where  $h$  is the hash commitment in  $\pi$
2. **Temporal Ordering:** Commitment  $C$  finalized at consensus time  $t$  with finality guarantees
3. **Authorship Attribution:** Commitment signed by holder of private key corresponding to  $pk$

**Critical constraint:** Verification must not require trust in RPC endpoints, platform APIs, or data availability committees beyond the Byzantine consensus assumption  $f < n/3$ .

### 3. Formal Verification Requirements

We formalize properties necessary for independent verification under the adversarial model of Section 2.

#### 3.1 Verification Independence

**Definition 1 (Independent Verification).** A verification system achieves *independence* if the predicate  $\text{Verify}(m, \pi, H) \rightarrow \{\text{true}, \text{false}\}$  is computable by any party possessing  $(m, \pi, H)$  without querying trusted external authorities beyond peer-to-peer consensus header retrieval.

This definition excludes systems requiring: (a) queries to trusted full nodes maintaining global state, (b) certificate authority validation, (c) platform API access, or (d) centralized timestamping services.

#### 3.2 Resource Feasibility Constraints

**Definition 2 (Browser-Native Verification).** Verification satisfies *browser-native* constraints if resource consumption bounds are:

- Computation:  $O(1)$  cryptographic operations +  $O(|m|)$  hashing
- Memory: Working set  $\leq 1$  MB
- Bandwidth: Proof size  $|\pi| \leq 1$  KB
- Latency: Total verification  $\leq 100$  ms (excluding artifact transfer)

These bounds derive from empirical constraints of mobile browsers and represent conservative thresholds for user-acceptable performance.

#### 3.3 Account-Local Verification

**Definition 3 (Local Verification Property).** A system exhibits *account-local verification* if validating commitment C for account A requires processing only state data dependent on A, independent of accounts  $B_1, \dots, B_k$  for  $k \rightarrow \infty$ .

Formally: verification complexity must be  $O(1)$  with respect to total account count N.

This constraint eliminates architectures requiring global state reconstruction or processing transactions from unrelated accounts during verification.

#### 3.4 Proof-First Consensus

**Definition 4 (Commitment-Based Consensus).** A consensus mechanism is *proof-first* if it achieves Byzantine agreement on commitments to state descriptors rather than execution correctness.

Specifically, consensus produces ordered commitment tuples:

$$C = (h, t, pk) \text{ where:}$$

```
h ∈ {0,1}^λ [content hash]  
t ∈ ■ [consensus timestamp]  
pk ∈ PK [public key]
```

The consensus layer validates structural properties (signature validity, hash chain integrity) but does not execute or validate semantic correctness. This separation provides two benefits:

1. **Reduced consensus complexity:** No smart contract execution during ordering phase
2. **Protocol neutrality:** Content interpretation remains external to consensus mechanism

### 3.5 Censorship-Resistant Verification

**Definition 5 (Verification Persistence).** A system has *persistent verification* if  $\text{Verify}(m, \pi, H) \rightarrow \text{true}$  remains evaluable after adversary ■ removes artifact  $m$  from all centralized platforms under its control.

This requires proof bundles  $\pi$  and consensus headers  $H$  remain accessible through decentralized peer-to-peer mechanisms. Note this concerns verification capability persistence, not content availability—archived copies of  $m$  may become inaccessible while verification remains possible if alternative copies emerge.

## 4. Verification Complexity and Proof Size Analysis

We establish quantitative bounds demonstrating browser-native verification feasibility under stated cryptographic assumptions.

### 4.1 Proof Bundle Size Bounds

**Lemma 1 (Minimal Proof Size).** Under cryptographic assumptions of Section 2.2, minimal proof bundles satisfy  $128 \leq |\pi| \leq 192$  bytes depending on signature scheme selection.

*Proof.* A proof bundle  $\pi$  contains:

- Content hash:  $h = H(m)$  with  $|h| = 32$  bytes (SHA-256,  $\lambda = 256$ )
- Digital signature:  $\sigma = \text{Sign}_{\text{sk}}(h)$  with  $|\sigma| \in \{64, 96\}$  bytes (Ed25519 vs BLS12-381)
- Consensus reference: Header identifier and Merkle authentication path, 24–32 bytes
- Temporal metadata: Timestamp  $t$ ,  $|t| = 8$  bytes (64-bit Unix time)
- Account identifier:  $|id| \in \{8, 32\}$  bytes depending on addressing scheme

Summing components:  $32 + 64 + 24 + 8 + 8 = 136$  bytes (Ed25519, minimal) to  $32 + 96 + 32 + 8 + 32 = 200$  bytes (BLS, maximal metadata). Conservative bound:  $128 \leq |\pi| \leq 192$ . ■

This represents 1–2 orders of magnitude improvement over Merkle inclusion proofs in large state trees ( $O(\log N) \times 32$  bytes  $\approx 640$  bytes for  $N = 2^{22}$  accounts).

### 4.2 Verification Complexity Bounds

**Lemma 2 (Constant Verification Complexity).** For dual-ledger architectures satisfying Definition 3, verification complexity is  $O(1)$  with respect to total account count  $N$ .

*Proof.* Verification algorithm executes:

1. Hash computation:  $h' = H(m)$ , time  $O(|m|)$ , artifact-dependent only
2. Signature verification:  $\text{Verify}_{\text{pk}}(\sigma, h)$ ,  $O(1)$  operation (~0.1 ms for Ed25519)
3. Consensus finality check: Validate quorum signatures,  $O(k)$  where  $k = \lceil 2n/3 \rceil$  is constant
4. Merkle authentication: Verify path in momentum,  $O(\log d)$  where  $d = \text{tree depth} \ll \log N$

Critical observation: Step 1 depends on  $|m|$  (artifact size), not  $N$  (account count). Steps 2–4 are  $O(1)$  in  $N$  because account-local property (Definition 3) eliminates dependency on global account population. Therefore, total complexity  $O(1)$  in  $N$ . ■

Operations explicitly *not* required: blockchain history download, unrelated transaction processing, smart contract execution, global state maintenance.

### 4.3 Empirical Performance Estimates

**Remark (Implementation Benchmarks).** Based on standard cryptographic library performance:

- SHA-256 hashing: ~400 MB/s on commodity hardware (~2.5 ms for 1 MB document)
- Ed25519 signature verification: ~70,000 operations/second (~0.014 ms per signature)
- Merkle path verification: ~100,000 operations/second (~0.01 ms per hash)

Estimated total verification latency: <3 ms for typical documents, dominated by artifact hashing in Step 1. This satisfies browser-native latency constraint (<100 ms, Definition 2).

**Caveat:** These are implementation-dependent estimates based on reference implementations, not architectural guarantees. The architectural claim is O(1) complexity in N; absolute timing varies with hardware and implementation quality.

#### 4.4 Architectural Comparison

We compare verification requirements across architecture classes under identical cryptographic assumptions (SHA-256, Ed25519,  $f < n/3$  Byzantine tolerance).

<b>Architecture</b>	<b>Proof Size</b>	<b>Complexity in N</b>	<b>Trust Dependencies</b>
Centralized Platform	N/A (API-based)	N/A	Platform operator
SPV (Bitcoin-like)	$O(\log N) \sim 640B$	$O(\log N)$	Full node for headers
Ethereum Light Client	$\sim 1 \text{ KB} + \text{state}$	$O(\log N) + \text{RPC}$	RPC endpoint + state
Optimistic Rollup	$\sim 500 \text{ bytes} + \text{calldata}$	$O(\log N) + \text{sequencer}$	Sequencer + DA layer
Dual-Ledger Header-First	128–192 bytes	$O(1)$	Consensus ( $f < n/3$ )

**Observation:** Only header-first architectures with account-isolated state achieve simultaneous  $O(1)$  verification complexity and minimal trust assumptions (consensus only). Alternative architectures optimize for different objectives (throughput, programmability) accepting higher verification costs or additional trust dependencies.

## 5. Comparative Architectural Analysis

We analyze why established architectural patterns do not satisfy all requirements (Definitions 1–5) concurrently. This analysis is structural and non-evaluative—each architecture optimizes for specific design objectives.

### 5.1 Centralized Trust Architectures

**Constraint violations:** Definitions 1 (independence), 5 (persistence).

Centralized platforms—social networks, cloud infrastructure, PKI systems—provide verification only through operator-controlled APIs. Structural implications:

- Verification requires platform infrastructure queries (violates Definition 1)
- Platform compromise eliminates verification for all dependent parties
- Content removal prevents verification by design (violates Definition 5)

These systems prioritize operational efficiency and user experience over adversarial robustness. Under threat model Section 2.1, they constitute single points of failure.

### 5.2 Global State Consensus Systems

**Constraint violations:** Definitions 2 (resource feasibility), 3 (local verification).

Account-based blockchains (Ethereum, similar designs) maintain global state requiring validators to process all transactions. Verification implications:

1. **Light client limitations:** Verification without full state requires trusted RPC endpoints (practical violation of Definition 1)
2. **Execution coupling:** Verification requires replaying smart contract execution, not merely validating cryptographic proofs
3. **State bloat:** Full verification requires synchronizing gigabytes of historical state (violates Definition 2 resource bounds)

These systems prioritize programmability and composable state over verification independence. The design trade-offs are intentional for their use cases.

### 5.3 Layer-2 Scaling Architectures

**Constraint violations:** Definition 1 (introduces additional trust dependencies).

Rollup architectures improve base layer throughput through off-chain execution. Trust implications:

- **Sequencer dependency:** Centralized operator controls transaction ordering (trust assumption even in optimistic designs with fraud proofs)
- **Data availability requirements:** Requires external committees or base layer calldata posting

- **Verification path complexity:** Validity proofs still require base layer verification chain

These mechanisms optimize for transaction throughput rather than verification minimalism. The additional trust assumptions are acceptable for their design objectives but do not satisfy Definition 1 independence constraint.

## 5.4 Content-Addressed Storage

**Constraint violations:** Definition 4 (lacks temporal ordering and authorship).

Content-addressed systems (IPFS, related approaches) provide integrity through cryptographic hashing but lack:

- Byzantine consensus for temporal ordering (when did content exist?)
- Cryptographic authorship attribution (who created it?)
- Availability guarantees under adversarial conditions

Hybrid approaches combining content addressing with blockchain timestamping reintroduce constraints from Section 5.2. These systems optimize for data deduplication and peer-to-peer distribution.

## 5.5 Summary of Comparative Analysis

Existing architectures optimize for legitimate objectives: operational efficiency (centralized), programmability (global state), throughput (rollups), or distribution (content addressing). No existing architecture simultaneously satisfies all five definitions without introducing additional trust assumptions or violating resource constraints. This motivates investigating alternative architectural approaches in Section 6.

## 6. Necessary Architectural Properties

Having established which architectures do not satisfy all constraints, we derive minimal properties required for concurrent satisfaction of Definitions 1–5.

### 6.1 Account State Isolation

**Property 1 (State Independence).** Verification of commitment C for account A must not require processing state from unrelated accounts B $\blacksquare$ , ..., B $\blacksquare$ .

Dual-ledger architectures achieve this through:

1. **Account chains:** Per-account transaction sequences maintaining local state evolution
2. **Consensus layer:** Global ordering mechanism producing commitments without executing state transitions

This separation enables verifying "account A committed hash h at time t" by examining: (1) A's local chain, (2) consensus layer inclusion proof—without processing accounts B through Z.

**Theorem 1 (Necessity of State Isolation).** For verification complexity  $O(1)$  in account count N under resource constraints Definition 2, account state must be isolated from global execution.

*Proof.* Suppose verification of account A requires processing  $k(N)$  accounts where  $k(N)$  grows with N. Then verification cost is  $\Omega(k(N))$ . Browser-native constraint (Definition 2) requires bounded computation, implying  $k(N) = O(1)$ . Therefore, state must be isolated. ■

### 6.2 Header-First Consensus Protocol

**Property 2 (Commitment Finalization).** Consensus must achieve Byzantine agreement on commitments  $C = (h, t, pk)$ , not execution outcomes.

The consensus layer produces header sequence  $H\blacksquare, H\blacksquare, \dots, H\blacksquare$  where header  $H\blacksquare$  contains:

- Merkle root  $R\blacksquare$  of account chain references
- Consensus timestamp  $t\blacksquare$
- Previous header hash:  $h\blacksquare = H(H\blacksquare\blacksquare\blacksquare)$
- Byzantine quorum signatures  $S\blacksquare$  where  $|S\blacksquare| \geq 2n/3$

Verification algorithm: (1) validate signature quorum, (2) verify hash chain integrity  $H(H\blacksquare\blacksquare\blacksquare) = h\blacksquare$ , (3) check Merkle inclusion in  $R\blacksquare$ . Transaction execution is not required.

### 6.3 Bounded Proof Construction

**Property 3 (Constant Proof Size).** Proof bundle size  $|\pi| = O(1)$ , specifically  $|\pi| < 1$  KB (Definition 2 constraint).

Achieved through: (a) direct header references eliminating deep Merkle trees, (b) deterministic commitment structure requiring no auxiliary data, (c) efficient signature schemes (64–96 bytes).

Lemma 1 establishes 128–192 byte bounds, significantly below 1 KB threshold with safety margin.

#### **6.4 Protocol-Level Content Neutrality**

**Property 4 (Semantic Independence).** Protocol-layer verification must not enforce content policies or semantic validation.

Rationale: Under adversarial conditions (Section 2.1), content filtering at protocol layer enables censorship vectors. Proof-first consensus (Property 2) naturally provides this—consensus validates commitment existence and structural properties, not content semantics.

Content moderation, if desired, occurs at application layer through reputation systems, user-side filtering, or client policies—not protocol enforcement mechanisms.

## 7. Architectural Case Study Analysis

We analyze reference implementations exhibiting Properties 1–4, evaluating architectural compliance without endorsing specific systems.

### 7.1 Dual-Ledger Architecture Description

Implementations separate account chains from consensus momentums:

- **Account chains:** Each account maintains hash-linked transaction sequence. Transactions reference previous account states exclusively, not global state. State transitions are deterministic given account chain history.
- **Momentum layer:** Byzantine consensus produces ordered commitments termed "momentums." Each momentum contains Merkle root of account chain updates finalized in that consensus round. The layer validates commitment structure without executing transactions.

### 7.2 Verification Protocol Specification

Verification of artifact  $m$  with proof bundle  $\pi$  proceeds as follows:

1. Compute  $h = H(m)$  using agreed hash function
2. Extract signature  $\sigma$  and public key  $pk$  from  $\pi$
3. Verify signature:  $\text{Verify\_pk}(\sigma, h) \rightarrow \{\text{true}, \text{false}\}$
4. Retrieve momentum header referenced in  $\pi$  via peer-to-peer protocol
5. Verify Merkle inclusion: account chain commitment  $\in$  momentum Merkle root
6. Validate momentum signatures: check  $|\{\text{valid signatures}\}| \geq \lceil 2n/3 \rceil$
7. Extract timestamp  $t$  from validated momentum header

**Complexity analysis:** One hash  $O(|m|)$ , one signature verification  $O(1)$ , one Merkle path verification  $O(\log d)$  where  $d$  = momentum tree depth (typically  $d < 20$ ), one quorum check  $O(n)$ . All operations independent of total account count  $N$ .

### 7.3 Compliance Verification

Formal verification against Section 3 definitions:

- **Independence (Def. 1):** Satisfied. Verification requires only  $(m, \pi, \text{headers})$  obtainable via peer-to-peer protocol, no RPC queries to trusted nodes
- **Browser-Native (Def. 2):** Satisfied. Proof size 128–192 bytes (Lemma 1), complexity  $O(1)$  in  $N$  (Lemma 2), empirical latency  $< 3$  ms
- **Account-Local (Def. 3):** Satisfied. Verification complexity  $O(1)$  in account count by construction satisfying Theorem 1 requirement
- **Proof-First (Def. 4):** Satisfied. Consensus achieves agreement on commitment tuples  $(h, t, pk)$  without validating execution

- **Persistence (Def. 5):** Satisfied. Headers and account chains distributed via gossip protocol independent of centralized infrastructure

**Architectural conclusion:** Dual-ledger architectures with header-first consensus satisfy all five definitions concurrently. This represents the only known architecture class achieving this property under stated resource and trust constraints.

## 8. Application Domains

We demonstrate practical utility through concrete verification scenarios under the established architectural framework.

### 8.1 Tamper-Evident Document Publication

**Scenario:** Journalist publishes investigative report requiring proof against retroactive alteration or suppression.

**Protocol execution:**

1. Journalist computes  $h = H(\text{report})$ , generates signature  $\sigma = \text{Sign}_{\text{sk}}(h)$
2. Commitment  $C = (h, t, \text{pk})$  finalized in consensus momentum
3. Report published with proof bundle  $\pi = (h, \sigma, \text{momentum\_ref}, t)$

**Verification:** Any party independently validates  $H(\text{downloaded\_report}) = h$  and verifies momentum inclusion. Functions even after journalist's website is compromised or removed.

**Explicit limitation:** System proves document existed at time  $t$  in form  $h$ —does not validate content truthfulness. Semantic validation remains recipient responsibility.

### 8.2 Pseudonymous Source Continuity

**Scenario:** Anonymous source publishes multiple leaks over time. Recipients must verify all originate from identical source without identity disclosure.

**Protocol execution:**

1. Source generates keypair  $(\text{pk}, \text{sk})$ , publishes  $\text{pk}$  via secure channel
2. Each leak signed with  $\text{sk}$ , commitment finalized in consensus
3. Consensus provides temporal ordering via momentum timestamps

**Verification:** Recipients verify all leaks share public key  $\text{pk}$ , establishing source continuity. Under key compromise, source rotates keys via signed transition statement committed to consensus.

**Explicit limitation:** System provides cryptographic key continuity, not real-world identity binding. Establishing source credibility requires external reputation mechanisms.

### 8.3 Capabilities and Fundamental Limitations

The architecture provides *integrity and temporal verification exclusively*. It explicitly does not provide:

- Content truthfulness validation (requires external fact-checking mechanisms)
- False information prevention (content neutrality is architectural property)
- Content availability guarantees (censorship resistance of distribution is orthogonal)

- Real-world identity binding (only cryptographic key continuity)
- Credibility or reputation assessment (application-layer concern)

These limitations are fundamental to the architecture—verification infrastructure provides tools for establishing provenance, not arbitration of truth. The critical distinction between proving *what was claimed* versus proving *what is true* must remain unambiguous.

## 9. Ethical Considerations

### 9.1 Misinformation and False Content

**Concern:** Malicious actors can timestamp false information, potentially lending false credibility.

**Analysis:** The system proves "entity X claimed Y at time T," not "Y is true." This is functionally equivalent to notarization—proving signature occurrence without validating content correctness. Importantly, false claims exist under all information architectures. Verification systems make provenance transparent, enabling repeated false claims from public key pk to accumulate verifiable history—a capability unavailable in opaque centralized platforms where attribution can be retroactively altered.

### 9.2 Coordinated Disinformation

**Concern:** Infrastructure could facilitate coordinated disinformation campaigns.

**Analysis:** Coordinated disinformation exists on centralized platforms, often amplified through algorithmic recommendation systems. Verification systems are content-neutral by design and lack amplification mechanisms. Transparent provenance enables researchers, journalists, and civil society organizations to track disinformation networks more effectively than on opaque platforms where coordination can occur without detection.

### 9.3 Immutability and Data Deletion

**Concern:** Commitments are immutable, potentially conflicting with data deletion requirements.

**Analysis:** Immutability is fundamental to tamper-evident architectures—permitting retroactive modification defeats integrity guarantees. However, immutability applies to *commitment hashes*, not content distribution. Content remains removable from hosting platforms, search indices, and archives. Hash commitment  $h = H(m)$  persists, but without  $m$ , provides no information beyond "something with hash  $h$  existed at time  $t$ ." This satisfies practical privacy objectives while preserving integrity verification for parties possessing content.

### 9.4 Neutral Infrastructure Trade-offs

We argue neutral verification infrastructure presents preferable failure modes compared to centralized gatekeeping:

1. **Cryptographic accountability:** Public key histories enable reputation assessment without centralized arbiters
2. **Elimination of single points of failure:** Centralized platforms can be compromised, coerced, or corrupted—distributed verification removes this vulnerability
3. **Transparent provenance:** Neutral systems expose information flows; centralized platforms can manipulate visibility and ordering without detection
4. **Evidence preservation:** For human rights documentation, legal proceedings, and investigative journalism, tamper-evident records protect evidence from suppression

While neutral systems can be misused, transparency enables countermeasures unavailable in opaque centralized platforms. The failure mode of neutrality—transparent misuse subject to analysis and reputation consequence—is preferable to centralized control failure mode—opaque manipulation without detection or accountability.

## 10. Open Problems and Research Directions

### 10.1 Cryptographic Proof Aggregation

**Problem statement:** Verifying  $N$  artifacts requires  $N$  independent proofs, total size  $O(N)$ .

**Research direction:** Investigate zero-knowledge proof systems (SNARKs, STARKs) for aggregating multiple verification proofs into constant-size aggregate proofs. Primary challenge: maintaining browser-native verification given zero-knowledge proof verification complexity.

### 10.2 Cross-Architecture Verification Standards

**Problem statement:** Different verification systems employ incompatible proof formats.

**Research direction:** Develop standardized verification interfaces enabling interoperability between proof-first architectures. Analogous to W3C DID standards for decentralized identity. Would enable universal browser extensions for verification across heterogeneous systems.

### 10.3 Eclipse Attack Resistance Analysis

**Problem statement:** Light clients trusting small peer sets can be eclipsed by adversarial nodes.

**Research direction:** Formal game-theoretic analysis of peer sampling strategies providing probabilistic eclipse resistance guarantees. Requires modeling adversarial network participation economics and information-theoretic peer diversity bounds.

### 10.4 Privacy-Preserving Verification

**Problem statement:** Public key exposure enables social graph analysis through commitment tracking.

**Research direction:** Integrate ring signatures or group signatures enabling verification without revealing exact signing key. Primary challenge: analyzing trade-off between proof size increase and browser-native compatibility constraints.

### 10.5 Formal Implementation Verification

**Problem statement:** Architectural analysis does not guarantee implementation correctness.

**Research direction:** Apply formal methods (Coq, TLA+, Isabelle/HOL) to verify implementations correctly realize specified properties. Essential for security-critical applications where implementation bugs could compromise integrity guarantees.

## 11. Conclusion

### 11.1 Summary of Contributions

This work establishes:

1. Formal definitions of independent verification with explicit resource bounds (Definitions 1–5)
2. Byzantine adversarial model with explicit fault tolerance thresholds ( $f < n/3$ ), consensus safety/liveness separation, and out-of-scope attack boundaries
3. Complexity bounds with formal justification: 128–192 byte proofs (Lemma 1),  $O(1)$  verification in account count (Lemma 2 and Theorem 1), empirical performance estimates
4. Systematic comparative analysis demonstrating existing architectures do not satisfy all constraints concurrently due to fundamental design trade-offs
5. Derivation of necessary architectural properties: account state isolation (Theorem 1), header-first consensus, bounded proof construction, protocol content neutrality
6. Architectural compliance verification demonstrating dual-ledger implementations satisfy all requirements

### 11.2 Formally Established Claims

Claims with rigorous justification:

- Browser-native verification requires  $O(1)$  complexity in global accounts (Theorem 1)
- Minimal proof bundles bounded by 128–192 bytes under stated cryptographic assumptions (Lemma 1)
- Dual-ledger verification achieves  $O(1)$  complexity in account count (Lemma 2)
- Existing architectures introduce additional trust assumptions or violate resource constraints (Section 5 structural analysis)

### 11.3 Open Conjectures

Hypotheses requiring additional investigation:

- Zero-knowledge proof aggregation can achieve  $O(1)$  proof size for  $N$  artifacts while maintaining browser-compatible verification complexity
- Eclipse attack resistance is achievable with bounded peer sampling under specific adversarial economic models
- Privacy-preserving verification through ring/group signatures is compatible with browser-native resource constraints

### 11.4 Architectural Implications

This work establishes that independent media verification is architecturally feasible without trusted third parties under Byzantine adversarial conditions. We have identified the minimal set of properties required and demonstrated they are realizable through dual-ledger designs with

header-first consensus.

This architectural capability does not address orthogonal problems: misinformation detection, content moderation policy, or reputation mechanism design. What this architecture provides is a neutral substrate—a content-agnostic layer enabling cryptographic verification of artifact integrity and temporal ordering independent of centralized authorities.

The fundamental distinction between proving *what was claimed* (within scope) versus proving *what is true* (outside scope) is central to understanding this contribution. We address only the former—establishing provenance without judging content validity.

Whether this capability improves information ecosystems depends on deployment context, complementary mechanisms (reputation systems, fact-checking infrastructure), and social adaptation patterns—determinations requiring empirical study beyond architectural analysis scope.

We conclude that dual-ledger architectures with proof-first consensus represent a necessary evolution in verification infrastructure design, providing capabilities unattainable in existing architectures while maintaining rigorous security properties under Byzantine fault tolerance assumptions.